

Introduction to Python

Basics of Data Analysis Using Pandas

David Durden, Data Services Librarian

Yishan Ding, Statistical Consulting Graduate Assistant



What is Python?

- An interpreted high-level general-purpose programming language
- Both a programming language and a tool that executes scripts written in Python
- Python was created Guido van Rossum and released in 1991

Why use Python?

- Free
- Open-source with active development
- Available on all major OS's (MacOS, Linux, Windows)
- Large community
- Easy-to-learn*
- Encourages re-use and reproducibility
- Interdisciplinary and extensible



Species of Python

Anaconda (<https://www.anaconda.com/>) is a specific distribution of Python that includes a lot of popular packages by default, including IPython console, Jupyter Notebook, Pandas, and others.

Several OS's come with Python installed by default. However, the system Python is often used by the operating system. Modifying the system Python *can* result in some programs not working correctly.

Outline for Today's Workshop

- Some basic Python commands
- Introduction to Pandas for data analysis
- Data import and preview
- Subsetting data
- Summary statistics and reporting
- Simple plots
- Simple Linear Regression

Set Up Your Environment

We will be using a Jupyter Notebook for this workshop to avoid having to edit text files and execute them at the Command Line.

Python can be run in interactive mode: You can type commands directly into the Python interpreter and see results. Open your Command Prompt or Terminal application and type `python`.

Using the interpreter is useful for testing code and getting immediate results. However, you cannot easily edit those commands or save your work.

Creating Python Scripts

Instead of using the interpreter, we can also write Python programs and then run them.

To get started, use a plain text editor like:

- BEdit (<https://www.barebones.com/products/bbedit/index.html>)
- Atom (<https://atom.io/>)
- Notepad++ (<https://notepad-plus-plus.org/>)

Text documents should be set up a specific way in order to read a valid Python code.

- The first line of your document should always be: `#!/usr/bin/env python`
- Save your file with `.py`

When you have written a program, you can run it using the following basic command:

```
python my_program.py
```

Getting Started

- Meet the print statement
 - `print("Hello World!")`

```
In [1]: print("Hello World!")
```

```
Hello World!
```

Fun fact: Printing *Hello world!* to the screen is a rhetorical function used by programmers to test that their systems are operating correctly. Incidentally, this has also become one of the first programs that you learn in almost every programming language. The creation of this function is attributed to Brian Kernighan, who first published it in a code snippet in *A Tutorial Introduction to the Programming Language B* in 1973.

Assign values to a variable

You can store anything in a variable in Python.

- `x = 3` # Assign a single integer value
- `y = [1, 2, 3]` # Assign a list of integer values
- `z = ["a", "b", "c"]` # Assign a list of character values
- `a = [y, z]` # Assign a list of variables
- `b = a` # Assign a variable to a variable

TIP: You can either 'call' a variable directly or print it using the `print()` function.

Variables

```
In [32]: x = 3 # Assign a single integer value
        y = [1, 2, 3] # Assign a list of integer values
        z = ["a", "b", "c"] # Assign a list of character values
        a = [x, y] # Assign a list of variables
        b = a

        print(x)
        print(y)
        print(z)
        print(a)
        print(b)

3
[1, 2, 3]
['a', 'b', 'c']
[3, [1, 2, 3]]
[3, [1, 2, 3]]
```

Using Python as a Calculator

- $35 * 40 / 4$
- $((6 - 4) * (2 + 5)) / 5$
- Advanced mathematical concepts are implemented in the `math` library which needs to be imported before we can use it
 - `from math import sqrt`
 - `sqrt(2)`

Do some Arithmetic

```
In [3]: 35*40/4
```

```
Out[3]: 350.0
```

```
In [4]: from math import sqrt  
sqrt(2)
```

```
Out[4]: 1.4142135623730951
```

5 Basic Data Types

| Type | Example |
|-----------|--|
| Integer | 3 or <code>int(3)</code> |
| Float | 3.1415 or <code>float(6/5)</code> |
| Character | "a", "b", "University of Maryland" or <code>str()</code> |
| Logical | True, False or <code>bool()</code> |
| Complex | 1+4j or <code>complex(1, 4)</code> |

- To determine an object type, use `type()`
- To test if an object is a specific type, use `isinstance()`
- To change type, use `int()`, `float()`, `str()`, `bool()`, `complex()`

Data Types

```
In [5]: type(3)
```

```
Out[5]: int
```

```
In [33]: type(3.1415)
```

```
Out[33]: float
```

```
In [34]: type("Testudo")
```

```
Out[34]: str
```

```
In [6]: foo = 6/5  
        isinstance(foo, float)
```

```
Out[6]: True
```

```
In [7]: complex(1, 4)
```

```
Out[7]: (1+4j)
```

Pandas: The Python Data Analysis Library

Pandas (<https://pandas.pydata.org/>) is already included by default in Anaconda Python. If you are using a different distribution you will have to download and install Pandas.

Pandas provides easy-to-use methods for working with data structures and performing data analysis using the Python language and environment.

Many of the commands and capabilities are similar to R.

We have to import Pandas before we can use it, similar to how we imported the `math` library.

- *`import pandas`*
- *`import pandas as pd`*

Import Pandas

```
In [22]: # To get started using Pandas we have to import it first  
import pandas  
  
# Imported libraries can be aliased to make code shorter  
# and easier to type  
# pd is the common alias for Pandas  
import pandas as pd  
# This is an optional parameter to limit the amount of  
# rows printed to the screen  
pd.set_option('display.max_rows', 10)
```

Basic Commands

Here are the basic commands when working with DataFrames in Pandas

| | |
|---|---|
| <code>df.max()</code> <code>df.min()</code> <code>df.mean()</code> <code>df.median()</code> <code>df.sum()</code> <code>df.cumsum()</code> | <code>df.var()</code> <code>df.std()</code> <code>len(df.index) or df.count()</code> <code>df.describe()</code> <code>df[]</code> |
|---|---|

Import a Dataset

We will begin by importing the `GSSsubset.csv` file

Pandas uses the `read_csv()` function to import csv files.

File names are string data and should be wrapped in quotes!

```
In [23]: # Note that pd.read_csv is used because we imported pandas as pd
pd.read_csv("data/GSSsubset.csv")
# For Windows users:
#pd.read_csv("data\GSSsubset.csv")
```

Out[23]:

| | id | sex | degree | income | marital | age | height | weig |
|-----|------|--------|-------------|----------|---------------|-----|--------|------|
| 0 | 1 | MALE | BACHELOR | 60967.50 | DIVORCED | 53 | 72 | 190 |
| 1 | 2 | FEMALE | BACHELOR | 60967.50 | MARRIED | 26 | 60 | 97 |
| 2 | 4 | FEMALE | BACHELOR | 10161.25 | MARRIED | 56 | 68 | 160 |
| 3 | 14 | FEMALE | HIGH SCHOOL | 17551.25 | MARRIED | 40 | 65 | 156 |
| 4 | 16 | MALE | HIGH SCHOOL | 17551.25 | MARRIED | 56 | 66 | 210 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 989 | 2531 | MALE | HIGH SCHOOL | 1478.00 | NEVER MARRIED | 40 | 71 | 230 |
| 990 | 2535 | MALE | HIGH SCHOOL | 33255.00 | DIVORCED | 56 | 72 | 195 |
| 991 | 2536 | MALE | HIGH SCHOOL | 8313.75 | NEVER MARRIED | 24 | 68 | 145 |
| 992 | 2537 | MALE | HIGH SCHOOL | 27712.50 | NEVER MARRIED | 27 | 68 | 180 |

| | id | sex | degree | income | marital | age | height | weig |
|-----|------|--------|----------------|----------|---------|-----|--------|------|
| 993 | 2538 | FEMALE | HIGH SCHOOL | 15703.75 | WIDOWED | 71 | 63 | 140 |

994 rows × 9 columns

Let's Examine a Dataset

When importing a dataset we should assign it to a variable for future use.

Declare a variable name and assign the `read_csv()` function from earlier.

```
In [5]: # df is common short hand in Pandas for "DataFrame", which is the object type  
# we are creating when importing a csv file  
gss_df = pd.read_csv("data/GSSsubset.csv")  
# For Windows users:  
# gss_df = pd.read_csv("data\GSSsubset.csv")
```

Preview the First Few Lines of a Dataset

After we assign our dataset to a variable, we can preview it with `head()`

Syntax: `dataFrame.head()`

```
In [6]: gss_df.head()  
# head() can be used on any dataframe object!
```

Out[6]:

| | id | sex | degree | income | marital | age | height | weight | h |
|---|----|--------|----------------|----------|----------|-----|--------|--------|---|
| 0 | 1 | MALE | BACHELOR | 60967.50 | DIVORCED | 53 | 72 | 190 | 6 |
| 1 | 2 | FEMALE | BACHELOR | 60967.50 | MARRIED | 26 | 60 | 97 | 4 |
| 2 | 4 | FEMALE | BACHELOR | 10161.25 | MARRIED | 56 | 68 | 160 | 2 |
| 3 | 14 | FEMALE | HIGH SCHOOL | 17551.25 | MARRIED | 40 | 65 | 156 | 3 |
| 4 | 16 | MALE | HIGH SCHOOL | 17551.25 | MARRIED | 56 | 66 | 210 | 6 |

Examine a Dataset

`head()` let's us look at the first few lines of a file.

`tail()` will let us look that last few lines of a file.

`describe()` will give us summary information.

Preview with tail()

```
In [7]: gss_df.tail()
```

```
Out[7]:
```

| | id | sex | degree | income | marital | age | height | weight |
|-----|------|--------|-------------|----------|---------------|-----|--------|--------|
| 989 | 2531 | MALE | HIGH SCHOOL | 1478.00 | NEVER MARRIED | 40 | 71 | 230 |
| 990 | 2535 | MALE | HIGH SCHOOL | 33255.00 | DIVORCED | 56 | 72 | 195 |
| 991 | 2536 | MALE | HIGH SCHOOL | 8313.75 | NEVER MARRIED | 24 | 68 | 145 |
| 992 | 2537 | MALE | HIGH SCHOOL | 27712.50 | NEVER MARRIED | 27 | 68 | 180 |
| 993 | 2538 | FEMALE | HIGH SCHOOL | 15703.75 | WIDOWED | 71 | 63 | 140 |

Return Summary Statistics with describe()

```
In [8]: print(gss_df.describe())
```

| | id | income | age | height | weight | \ |
|-------|-------------|---------------|------------|------------|------------|---|
| count | 994.000000 | 994.000000 | 994.000000 | 994.000000 | 994.000000 | |
| mean | 1271.183099 | 36887.218352 | 44.488934 | 67.414487 | 181.318913 | |
| std | 734.068659 | 34702.256068 | 13.126865 | 4.037741 | 41.641615 | |
| min | 1.000000 | 369.500000 | 19.000000 | 57.000000 | 90.000000 | |
| 25% | 648.250000 | 15703.750000 | 33.000000 | 64.000000 | 150.000000 | |
| 50% | 1254.500000 | 27712.500000 | 44.000000 | 67.000000 | 175.000000 | |
| 75% | 1915.750000 | 49882.500000 | 55.000000 | 70.000000 | 205.000000 | |
| max | 2538.000000 | 158656.952000 | 79.000000 | 79.000000 | 410.000000 | |

| | hrswrk |
|-------|------------|
| count | 994.000000 |
| mean | 42.640845 |
| std | 14.394974 |
| min | 1.000000 |
| 25% | 38.000000 |
| 50% | 40.000000 |
| 75% | 50.000000 |
| max | 89.000000 |

Exercise 1: Import a Dataset

- Import the dataset `ISBN.csv`
- Preview the dataset
 - Call the variable that you assigned to `ISBN.csv`
 - Call the `head()` method on the DataFrame

Examine the Dataset

- `df.shape` where `df` is the name of your DataFrame
 - Returns row x column
- `df.columns`
 - Returns a list of the column names
- `df.index`
 - Returns row information (range and step)
- `df.info()`
 - Returns list of variables, data type, count, size, etc.

Exercise 2

- Examine the dimensions of `ISBN.csv`
 - How many rows does this dataset have?
 - How many columns?
 - What are the variable names?

```
In [24]: print(gss_df.shape)
```

```
(994, 9)
```

```
In [25]: print(gss_df.describe())
```

| | id | income | age | height | weight \ |
|-------|-------------|---------------|------------|------------|------------|
| count | 994.000000 | 994.000000 | 994.000000 | 994.000000 | 994.000000 |
| mean | 1271.183099 | 36887.218352 | 44.488934 | 67.414487 | 181.318913 |
| std | 734.068659 | 34702.256068 | 13.126865 | 4.037741 | 41.641615 |
| min | 1.000000 | 369.500000 | 19.000000 | 57.000000 | 90.000000 |
| 25% | 648.250000 | 15703.750000 | 33.000000 | 64.000000 | 150.000000 |
| 50% | 1254.500000 | 27712.500000 | 44.000000 | 67.000000 | 175.000000 |
| 75% | 1915.750000 | 49882.500000 | 55.000000 | 70.000000 | 205.000000 |
| max | 2538.000000 | 158656.952000 | 79.000000 | 79.000000 | 410.000000 |

| | hrswrk |
|-------|------------|
| count | 994.000000 |
| mean | 42.640845 |
| std | 14.394974 |
| min | 1.000000 |
| 25% | 38.000000 |
| 50% | 40.000000 |
| 75% | 50.000000 |
| max | 89.000000 |

```
In [40]: gss_df.columns
```

```
Index(['id', 'sex', 'degree', 'income', 'marital', 'age', 'height', 'weight',  
      'hrswrk'],  
      dtype='object')
```

```
In [96]: gss_df.index
```

```
Out[96]: RangeIndex(start=0, stop=994, step=1)
```

```
In [43]: gss_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 994 entries, 0 to 993  
Data columns (total 9 columns):  
id          994 non-null int64  
sex         994 non-null object  
degree      994 non-null object  
income      994 non-null float64  
marital     994 non-null object  
age         994 non-null int64  
height      994 non-null int64  
weight      994 non-null int64  
hrswrk      994 non-null int64  
dtypes: float64(1), int64(5), object(3)  
memory usage: 70.0+ KB
```

Subset the Dataset

| 1. Subset by row/column name | 2. Subset by row/column index |
|---|--|
| df.loc[rowname, colname] | df.iloc[rowindex, colindex] |
| - data.loc[2, 'degree'] - data.loc[:3, 'id': 'income'] - data.loc[[1,3,5], ['id', 'degree', 'income'] | - data.iloc[1, 3] - data.iloc[:5,] - data.iloc[900:, [2,3,5,8]] |

| 3. Subset by variable name | 4. Subset by specific criteria |
|---|---|
| df['colname'] df.colname | data[data['variable'] <operator> 'criteria'] |
| - data['id'] - data[['id', 'degree', 'income']] - data.degree | - data[data['age'] < 50] - data[data['sex'] == 'MALE'] - data[(data['age'] < 50) & (data['sex'] == 'MALE')] |

Refer to [this \(https://cmdlinetips.com/2019/03/how-to-select-one-or-more-columns-in-pandas/\)](https://cmdlinetips.com/2019/03/how-to-select-one-or-more-columns-in-pandas/) article for a more in-depth explanation of subsetting in Pandas.

```
In [26]: # Here is an example of a subset for males less than 50 years of age  
# Note that this expression includes the head() method!  
gss_df[(gss_df['age'] < 50) & (gss_df['sex'] == 'MALE')].head()
```

Out[26]:

| | id | sex | degree | income | marital | age | height | weight | h |
|----|----|------|----------------|----------|---------------|-----|--------|--------|---|
| 7 | 27 | MALE | HIGH SCHOOL | 17551.25 | SEPARATED | 35 | 70 | 180 | 4 |
| 11 | 38 | MALE | HIGH SCHOOL | 33255.00 | DIVORCED | 31 | 67 | 150 | 3 |
| 12 | 40 | MALE | BACHELOR | 73900.00 | MARRIED | 35 | 69 | 200 | 5 |
| 13 | 44 | MALE | HIGH SCHOOL | 24017.50 | NEVER MARRIED | 26 | 76 | 165 | 4 |
| 16 | 47 | MALE | JUNIOR COLLEGE | 40645.00 | NEVER MARRIED | 30 | 72 | 185 | 4 |

Exercise 3

Using the `ISBN.csv` dataset answer the following question:

- In which years was the *Journal of Crystal Growth* requested?

Hint: You will need to subset by criteria and then subset by column name!

- You can also sort data using the following code:
 - `df.sort_values(by=['colname'], inplace=False, ascending=False)`

Summary Report

- Summary of each variable: `df.describe()`
- Summary statistics for subgroups/categories:
 - **Example question:** According to the GSSsubset dataset, what is the average income by sex?

```
In [99]: gss_df.groupby(['sex'])[['income']].aggregate('mean')
```

Out[99]:

| | income |
|--------|--------------|
| sex | |
| FEMALE | 27300.446119 |
| MALE | 46095.814166 |

Exercise 4: Answering Questions with Data

- Import the dataset `lib_checkout.csv`
- How many books were checked out by undgrads in total?
- How many books were checked out by graduate students in total?
- How many books were checked out by day of the week on average?

Bonus question:

- *How many books were checked out by grad students and undergraduate students by day of the week on average? (use only **one** command)*
- *Assign the result to a variable named `weekmean`*

```
In [27]: # Read the lib_checkout dataset and store in a variable
libdata = pd.read_csv("data/lib_checkout.csv")
# For Windows users
#libdata = pd.read_csv("data\\lib_checkout.csv")

# Subset by undergrad
ulib = libdata[libdata['status'] == 'Undergrad']
```

```
In [11]: # Preview the subset  
ulib.head()
```

Out[11]:

| | date | day | count | status | laptop_checkout | book_checkout |
|---|----------|-----------|-------|-----------|-----------------|---------------|
| 0 | 9/1/2014 | Monday | 6570 | Undergrad | 329 | 1643 |
| 2 | 9/2/2014 | Tuesday | 5852 | Undergrad | 585 | 1522 |
| 4 | 9/3/2014 | Wednesday | 7060 | Undergrad | 847 | 1836 |
| 6 | 9/4/2014 | Thursday | 8065 | Undergrad | 968 | 2097 |
| 8 | 9/5/2014 | Friday | 6323 | Undergrad | 759 | 1770 |

Plotting in Python/Pandas

We need to import `matplotlib`, a popular scientific plotting library


```
In [12]: # This command forces plots to display in Jupyter Notebooks  
         %matplotlib inline  
         # This command imports matplotlib and assigns the alias plt  
         import matplotlib.pyplot as plt
```

Univariate Graph: Histogram

More information on plotting histograms [here \(https://pythonspot.com/matplotlib-histogram/\)](https://pythonspot.com/matplotlib-histogram/).

```
In [13]: # Graph the data
plt.hist(gss_df['income'], edgecolor='black', linewidth='1', alpha=0.5)
# Add title and axes labels
plt.title('Histogram of Income')
plt.xlabel('Annual Income')
plt.ylabel('Frequency')
```

```
Out[13]: Text(0,0.5,'Frequency')
```




Another way to Create a Histogram

The `plot()` function is useful for when you don't want to mess with the aesthetics of graph.

```
In [14]: gss_df['income'].plot(kind='hist', title='Histogram of Income').set(xlabel='Annual Income')
```


```
Out[14]: [Text(0.5,0,'Annual Income')]
```



Univariate Graph: Bar Chart

```
In [15]: # To plot a single variable on a bar chart, we need  
# to count the individual values first  
gss_df['marital'].value_counts().plot(kind='bar')
```


```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1188dc940>
```



Bivariate Graph: Scatterplot

```
In [16]: # Graph the data
plt.scatter(gss_df['hrswrk'], gss_df['income'])
# Set the axes and title labels
plt.title('Plot of Income and Working Hours')
plt.xlabel('Working Hours')
plt.ylabel('Annual Income')
```

```
Out[16]: Text(0,0.5,'Annual Income')
```



Exercise 5: Plot Some Data

Dataset: *lib_checkout.csv*

Draw a scatterplot showing `laptop_checkout` as the x variable and `book_checkout` as the y variable for **undergraduate students only**

Hint:

- *Create a subset of the original data by `status` and save to a new `DataFrame`*
- *Plot using the new `DataFrame`*

Simple Linear Regression

$$y = \beta_0 + \beta_1 x + e$$

- y : dependent variable
- x : independent variable
- e : error term
- β_0 : intercept
- β_1 : slope

Simple Linear Regression

1. **F-test p-value:** *Is this model statistically significant?*
2. **r^2 :** *How much variance in y can be explained by the variance in x ?*
3. **Intercept**
4. **Slope**

Linear Regressions Using `statsmodels.api`

```
In [17]: import numpy as np
import statsmodels.api as sm

# This is the difficult way of implementing a linear model in Python
y1 = gss_df['income'] # Dependent variable
x1 = gss_df['hrrwrk'] # Independent variable
y1, x1 = np.array(y1), np.array(x1) # Convert to numpy arrays
x1 = sm.add_constant(x1) # Add constant (column of 1's)
model = sm.OLS(y1, x1) # Define the model
fitted = model.fit() # Fit the model
print(fitted.summary()) # Print the results
```

/anaconda3/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: Future Warning: The pandas.core.datetools module is deprecated and will be removed in a future version. Please use the pandas.tseries module instead.

```
from pandas.core import datetools
```

OLS Regression Results

```
=====
=
Dep. Variable:          y    R-squared:          0.11
8
Model:                OLS    Adj. R-squared:      0.11
7
Method:               Least Squares    F-statistic:      132.
3
Date:                 Wed, 24 Jul 2019    Prob (F-statistic):    7.79e-2
9
Time:                 10:21:35    Log-Likelihood:      -1174
0.
No. Observations:      994    AIC:              2.348e+0
4
Df Residuals:          992    BIC:              2.349e+0
4
Df Model:              1
Covariance Type:      nonrobust
=====
```

Linear Regressions Using `statsmodels.formula.api`

```
In [18]: import statsmodels.formula.api as smf

# This is the easier way of implementing a linear model using
# R-like formulas
modell = smf.ols(formula='income ~ hrswrk', data=gss_df).fit() # Note the lowercase 'ols'
print(modell.summary()) # Required to use patsy formulae
```

OLS Regression Results

```
=====
Dep. Variable:          income    R-squared:                0.118
Model:                  OLS      Adj. R-squared:            0.117
Method:                 Least Squares    F-statistic:          132.3
Date:                  Wed, 24 Jul 2019    Prob (F-statistic):    7.79e-29
Time:                  10:22:05    Log-Likelihood:        -11740.
No. Observations:      994    AIC:                  2.348e+04
Df Residuals:          992    BIC:                  2.349e+04
Df Model:               1
Covariance Type:       nonrobust
=====
```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-----------|-----------|----------|--------|-------|-----------|----------|
| Intercept | 1621.4901 | 3235.485 | 0.501 | 0.616 | -4727.690 | 7970.670 |
| hrswrk | 827.0410 | 71.895 | 11.503 | 0.000 | 685.957 | 968.125 |

```
=====
Omnibus:                421.362    Durbin-Watson:           1.708
Prob(Omnibus):           0.000    Jarque-Bera (JB):        1713.357
Skew:                    2.027    Prob(JB):                0.00
Kurtosis:                7.993    Cond. No.                 141.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```


Exercise 6: Linear Regressions

Use the undergraduate subset created in Exercise 4

1. Run a simple linear regression of `book_checkout` (dependent variable, y) on `laptop_checkout` (independent variable, x)
2. Using the regression output, write down the regression model as comment in your code
3. Is the model significant on $\alpha = 0.05$?

```
In [29]: #y = ulib['book_checkout']
#x = ulib['laptop_checkout']
model = smf.ols(formula='book_checkout ~ laptop_checkout', data=ulib).fit()
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          book_checkout    R-squared:                0.216
Model:                  OLS              Adj. R-squared:           0.188
Method:                 Least Squares    F-statistic:             7.714
Date:                   Wed, 24 Jul 2019  Prob (F-statistic):      0.00967
Time:                   11:14:33          Log-Likelihood:          -217.97
No. Observations:       30              AIC:                    439.9
Df Residuals:           28              BIC:                    442.7
Df Model:                1
Covariance Type:        nonrobust
=====
```

```
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept      1316.3691    158.720      8.294      0.000      991.245    164
1.493
laptop_checkout  0.7280      0.262      2.777      0.010      0.191
1.265
=====
```

```
Omnibus:                1.573    Durbin-Watson:           2.523
Prob(Omnibus):          0.455    Jarque-Bera (JB):        1.193
Skew:                   0.481    Prob(JB):                0.551
Kurtosis:               2.832    Cond. No.                1.47e+03
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

[2] The condition number is large, $1.47e+03$. This might indicate that there are strong multicollinearity or other numerical problems.

Thank you!

David Durden

 urden@umd.edu (<mailto:urden@umd.edu>).

 [https://lib.umd.edu/data\(\)](https://lib.umd.edu/data()).

Yishan Ding

 ysding@umd.edu (<mailto:ysding@umd.edu>).

 [https://www.lib.umd.edu/rc/statistical-consulting\(\)](https://www.lib.umd.edu/rc/statistical-consulting()).